

```
import cv2
from google.colab.patches import cv2_imshow
image=cv2.imread('set image path')
cv2_imshow(image)
cv2.waitKey(0)
cv2.destroyAllWindows()
gray_image = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY)
cv2_imshow(gray_image)
cv2.waitKey(0)
cv2.destroyAllWindows()
resized_image = cv2.resize(image,(200,200))
cv2_imshow(resized_image)
cv2.waitKey(0)
cv2.destroyAllWindows()
blurred_image = cv2.GaussianBlur(image, (15, 15), 0)
cv2_imshow(blurred_image)
cv2.waitKey(0)
cv2.destroyAllWindows()
edges = cv2.Canny(gray_image, 100, 200)
cv2_imshow(edges)
cv2.waitKey(0)
cv2.destroyAllWindows()
cv2.rectangle(image, (50,50), (300,300), (255, 0, 0), 2)
cv2.line(image, (60,60), (300,300), (0, 0, 255), 2)
cv2_imshow(image)
cv2.waitKey(0)
cv2.destroyAllWindows()
```

```

from PIL import Image
import numpy as np

def shift_image(img, depth_img, shift_amount=10):
    # Ensure base image has alpha
    img = img.convert("RGBA")
    data = np.array(img)
    # Ensure depth image is grayscale (for single value)
    depth_img = depth_img.convert("L")
    depth_data = np.array(depth_img)
    deltas = ((depth_data / 255.0) * float(shift_amount)).astype(int)
    # This creates the transparent resulting image.
    # For now, we're dealing with pixel data.
    shifted_data = np.zeros_like(data)
    height, width, _ = data.shape
    for y, row in enumerate(deltas):
        for x, dx in enumerate(row):
            if x + dx < width and x + dx >= 0:
                shifted_data[y, x + dx] = data[y, x]
    # Convert the pixel data to an image.
    shifted_image = Image.fromarray(shifted_data.astype(np.uint8))
    return shifted_image

img = Image.open("C:\\Users\\student\\Desktop\\cube1.jpeg")
depth_img = Image.open("C:\\Users\\student\\Desktop\\cube3.jpeg")
shifted_img = shift_image(img, depth_img, shift_amount=10)
shifted_img.show()

```

```

import cv2

def detect_moving_objects(video_path):
    cap = cv2.VideoCapture(video_path)
    if not cap.isOpened():
        print("Error: Couldn't open the video file.")
        return
    bg_subtractor = cv2.createBackgroundSubtractorMOG2()
    while cap.isOpened():
        ret, frame = cap.read()
        if not ret:
            break
        fg_mask = cv2.medianBlur(bg_subtractor.apply(frame), 5)
        contours, _ = cv2.findContours(fg_mask, cv2.RETR_EXTERNAL,
cv2.CHAIN_APPROX_SIMPLE)
        for contour in contours:
            if cv2.contourArea(contour) > 100:
                M = cv2.moments(contour)
                if M["m00"] != 0:
                    cx = int(M["m10"] / M["m00"])
                    cy = int(M["m01"] / M["m00"])
                    # Draw a dot (small filled circle)
                    cv2.circle(frame, (cx, cy), 5, (255, 0, 0), 5)
        cv2.imshow('Moving Object Detection', frame)
        if cv2.waitKey(45) & 0xFF == ord('q'):
            break
    cap.release()
    cv2.destroyAllWindows()
    video_path = 'input.mp4'
    detect_moving_objects(video_path)

```

```

from transformers import VisionEncoderDecoderModel, ViTFeatureExtractor,
AutoTokenizer

import torch

from PIL import Image

import warnings

warnings.filterwarnings('ignore')

model = VisionEncoderDecoderModel.from_pretrained("nlpconnect/vit-gpt2-image-
captioning")

feature_extractor = ViTFeatureExtractor.from_pretrained("nlpconnect/vit-gpt2-image-
captioning")

tokenizer = AutoTokenizer.from_pretrained("nlpconnect/vit-gpt2-image-captioning")

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

model.to(device)

max_length = 16

num_beams = 4

gen_kwargs = {"max_length": max_length, "num_beams": num_beams}

def predict_step(image_paths):
    images = []
    for image_path in image_paths:
        i_image = Image.open(image_path)
        if i_image.mode != "RGB":
            i_image = i_image.convert(mode="RGB")
        images.append(i_image)
    pixel_values = feature_extractor(images=images, return_tensors="pt").pixel_values
    pixel_values = pixel_values.to(device)
    output_ids = model.generate(pixel_values, **gen_kwargs)
    preds = tokenizer.batch_decode(output_ids, skip_special_tokens=True)
    preds = [pred.strip() for pred in preds]
    return preds

predict_step(["C:\\Users\\student\\Downloads\\ss.jpg"])

```

```

from PIL import
Image import cv2
import numpy as np
import requests
image_url =
'sample.jpg'
image = Image.open(image_url)
image = image.resize((450, 250))
image.show()
cv2.waitKey(0)
cv2.destroyAllWindows()
image_arr =
np.array(image)
grey = cv2.cvtColor(image_arr, cv2.COLOR_BGR2GRAY)
blur = cv2.GaussianBlur(grey, (5, 5), 0)
dilated = cv2.dilate(blur, np.ones((3, 3)))
dilated =
cv2.dilate(blur, np.ones((3, 3)))
kernel = cv2.getStructuringElement(cv2.MORPH_ELLIPSE, (2, 2))
closing = cv2.morphologyEx(dilated, cv2.MORPH_CLOSE, kernel)
car_cascade_src = "vehicle.xml"
car_cascade = cv2.CascadeClassifier(car_cascade_src)
cars = car_cascade.detectMultiScale(closing, 1.1, 1)
cnt = 0

for (x, y, w, h) in cars:

    cv2.rectangle(image_arr, (x, y), (x + w, y + h), (255, 0, 0), 2)

    cnt += 1

annotated_image = Image.fromarray(image_arr)
annotated_image.show()
cv2.waitKey(0)
cv2.destroyAllWindows()

```

```

import cv2
import numpy as np
import matplotlib.pyplot as plt#
Load the image
image = cv2.imread('path_to_your_image.jpg')
gray = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY)
# Apply GaussianBlur to reduce noise and improve contour detectionblurred =
cv2.GaussianBlur(gray, (5, 5), 0)
# Perform edge detection
edged = cv2.Canny(blurred, 50, 150)# Find
contours
contours, _ = cv2.findContours(edged, cv2.RETR_EXTERNAL,
cv2.CHAIN_APPROX_SIMPLE)
# Draw contours on the original image
contour_image = image.copy()
cv2.drawContours(contour_image, contours, -1, (0, 255, 0), 2)# Display
the results
plt.figure(figsize=(10, 10))
plt.subplot(1, 3, 1)
plt.title('Original Image')
plt.imshow(cv2.cvtColor(image, cv2.COLOR_BGR2RGB))
plt.subplot(1, 3, 2)
plt.title('Edge Detection')
plt.imshow(edged, cmap='gray')
plt.subplot(1, 3, 3) plt.title('Contours')
plt.imshow(cv2.cvtColor(contour_image, cv2.COLOR_BGR2RGB))plt.show()

```

```

import numpy as np
import matplotlib.pyplot as plt
from skimage.feature import canny
from skimage import data, segmentation, morphology, filters
from skimage.color import rgb2gray, label2rgb
import scipy.ndimage as nd
plt.rcParams["figure.figsize"] = (12, 8)
%matplotlib inline
rocket = data.rocket()
rocket_wh = rgb2gray(rocket)
edges = canny(rocket_wh)
plt.imshow(edges, interpolation='gaussian')
plt.title('Canny detector')
plt.show()
fill_im = nd.binary_fill_holes(edges)
plt.imshow(fill_im)
plt.title('Region Filling')
plt.show()
elevation_map = filters.sobel(rocket_wh)
plt.imshow(elevation_map)
plt.title('Elevation Map')
plt.show()
markers = np.zeros_like(rocket_wh)
markers[rocket_wh < 0.1171875] = 1
# 30/255
markers[rocket_wh > 0.5859375] = 2
# 150/255
plt.imshow(markers)
plt.title('Markers')
plt.show()
segments = segmentation.watershed(elevation_map, markers)
plt.imshow(segments)
plt.title('Watershed Segmentation')
plt.show()
segments_filled = nd.binary_fill_holes(segments - 1)
label_rock, _ = nd.label(segments_filled)
image_label_overlay = label2rgb(label_rock, image=rocket_wh)
fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(24, 16), sharey=True)
ax1.imshow(rocket_wh)
ax1.contour(segments_filled, [0.8], linewidths=1.8, colors='w')
ax2.imshow(image_label_overlay)
plt.show()

```

```

import cv2
import numpy as np

def detect_shapes(image_path):
    image = cv2.imread(image_path)
    if image is None:
        print(f"Error: Failed to load image from {image_path}.")
        return
    gray = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY)
    gray_blurred = cv2.medianBlur(gray, 5)
    edges = cv2.Canny(gray, 50, 150, apertureSize=3)
    lines = cv2.HoughLines(edges, 1, np.pi / 180, 150)
    if lines is not None:
        for rho, theta in lines[:, 0]:
            a = np.cos(theta)
            b = np.sin(theta)
            x0 = a * rho
            y0 = b * rho
            x1 = int(x0 + 1000 * (-b))
            y1 = int(y0 + 1000 * (a))
            x2 = int(x0 - 1000 * (-b))
            y2 = int(y0 - 1000 * (a))
            cv2.line(image, (x1, y1), (x2, y2), (0, 0, 255), 2)
    circles = cv2.HoughCircles(gray_blurred, cv2.HOUGH_GRADIENT, 1, 20, param1=50, param2=30,
                                minRadius=1, maxRadius=40)
    if circles is not None:
        circles = np.uint16(np.around(circles))
        for i in circles[0, :]:
            cv2.circle(image, (i[0], i[1]), i[2], (0, 255, 0), 2)
            cv2.circle(image, (i[0], i[1]), 2, (0, 0, 255), 3)
    contours, _ = cv2.findContours(edges, cv2.RETR_TREE, cv2.CHAIN_APPROX_SIMPLE)
    for contour in contours:
        epsilon = 0.02 * cv2.arcLength(contour, True)
        approx = cv2.approxPolyDP(contour, epsilon, True)
        if len(approx) == 3:
            cv2.drawContours(image, [approx], 0, (0, 255, 255), 2)
        elif len(approx) == 4:
            cv2.drawContours(image, [approx], 0, (255, 0, 0), 2)
        elif len(approx) > 4:
            area = cv2.contourArea(contour)
            if area > 100:
                cv2.drawContours(image, [approx], 0, (255, 255, 0), 2)
    cv2.imshow('Shape Detection', image)
    cv2.waitKey(0)
    cv2.destroyAllWindows()
    detect_shapes("C:\\Users\\student\\Downloads\\circle.png")

```



```

import cv2

import matplotlib.pyplot as plt

image_path = 'C:/Users/student/Downloads/perfect-family-photo-session-by-rebecca-danzenbaker.webp' # Replace with your image path

image = cv2.imread(image_path)

if image is None:
    print(f"Error: Could not load image from {image_path}")
else:
    gray = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY)
    face_cascade = cv2.CascadeClassifier(cv2.data.harcascades +
    'haarcascade_frontalface_default.xml')
    faces = face_cascade.detectMultiScale(gray, scaleFactor=1.1,
    minNeighbors=5, minSize=(30, 30))
    for (x, y, w, h) in faces:
        cv2.rectangle(image, (x, y), (x+w, y+h), (255, 0, 0), 2)
    image_rgb = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)
    plt.imshow(image_rgb)
    plt.axis('off')
    plt.show()

import cv2
import pytesseract

pytesseract.pytesseract.tesseract_cmd =
'c:\\Users\\online\\AppData\\Local\\Programs\\Tesseract-OCR\\tesseract.exe'

def extract_text_from_image(image_path):
    image = cv2.imread(image_path)
    gray = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY)
    text = pytesseract.image_to_string(gray)
    return text

image_path = './sample.jpeg'
extracted_text = extract_text_from_image(image_path)
print("Extracted Text:")
print(extracted_text)

```

```

import cv2

import numpy as np

def region_of_interest(img, vertices):

    mask = np.zeros_like(img)
    cv2.fillPoly(mask, vertices, 255)
    masked_image = cv2.bitwise_and(img, mask)
    return masked_image

def draw_lines(img, lines):

    if lines is not None:

        for line in lines:

            for x1, y1, x2, y2 in line:

                cv2.line(img, (x1, y1), (x2, y2), (0, 255, 0), 5)

def process_image(image):

    height, width = image.shape[:2] # Convert
    the image to grayscale

    gray = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY) # Apply Gaussian
    blur

    blur = cv2.GaussianBlur(gray, (5, 5), 0) # Apply Canny
    edge detector

    edges = cv2.Canny(blur, 50, 150) # Define the
    region of interest

    vertices = np.array([[(50, height), (width//2 - 50, height//2 + 50), (width//2 + 50, height//2 + 50), (width - 50,
height)], dtype=np.int32)

    masked_edges = region_of_interest(edges, vertices)

    lines = cv2.HoughLinesP(masked_edges, rho=1, theta=np.pi/180, threshold=50, minLineLength=50,
maxLineGap=200)

    # Draw the lines on the original image
    line_image =
    np.zeros_like(image)
    draw_lines(line_image, lines)

    result = cv2.addWeighted(image, 0.8, line_image, 1, 0)
    return result

def main():

    cap = cv2.VideoCapture("C:\\Users\\student.SCASA1\\Downloads\\travel_road.mp4")

    while(cap.isOpened()):

        ret, frame = cap.read()
        if ret:
            result = process_image(frame)
            cv2.imshow('Lane
Detection', result)
            if cv2.waitKey(1) & 0xFF == ord('q'):
                break
            else:

                break

        cap.release()

    cv2.destroyAllWindows()

if __name__ == '__main__':
    main()

```

```

import cv2
from pyzbar.pyzbar import decode
def decode_qr_from_image(image_path):
    frame = cv2.imread(image_path)
    if frame is None:
        print(f"Failed to load image from {image_path}")
        return
    gray_frame = cv2.cvtColor(frame, cv2.COLOR_BGR2GRAY)
    decoded_objects = decode(gray_frame)
    for obj in decoded_objects:
        qr_data = obj.data.decode('utf-8')
        print('Data:', qr_data)
        points = obj.polygon
        if len(points) > 4:
            hull = cv2.convexHull(np.array([point for point in points], dtype=np.float32))
            hull = list(map(tuple, np.squeeze(hull)))
        else:
            hull = points
        n = len(hull)
        for j in range(0, n):
            cv2.line(frame, hull[j], hull[(j + 1) % n], (255, 0, 0), 3)
    cv2.imshow('QR Code Scanner', frame)
    cv2.waitKey(0)
    cv2.destroyAllWindows()
def main():
    image_path = "C:\\Users\\student.SCASA1\\Downloads\\OIP.jpg"
    decode_qr_from_image(image_path)
if __name__ == "__main__":
    main

```

12

pip install deepface

import cv2

import matplotlib.pyplot as plt from

deepface import DeepFace

img = cv2.imread(r'/content/321.jpg')

plt.imshow(img[:, :, : -1]) plt.show()

result = DeepFace.analyze(img,

actions = ['emotion'])

print(result)

```

import cv2
import numpy as np
from google.colab.patches import
cv2_imshow def calculate_distance(bbox1,
bbox2):
    center1 = (bbox1[0] + bbox1[2] // 2, bbox1[1] + bbox1[3] //
2) center2 = (bbox2[0] + bbox2[2] // 2, bbox2[1] +
bbox2[3] // 2)
    distance = np.sqrt((center1[0] - center2[0])**2 + (center1[1] - center2[1])**2)
    return distance
def draw_bounding_box(image, bbox, color):
    cv2.rectangle(image, (bbox[0], bbox[1]), (bbox[0] + bbox[2], bbox[1] + bbox[3]), color,
2) image_path = '/content/ii.jpg'
    if not os.path.exists(image_path):
        print(f"Error: Image file '{image_path}' not
found.") else:
        image =
        cv2.imread(image_path) if
        image is None:
            print(f"Error: Unable to load image '{image_path}'")
        else:
            bbox1 = (100, 50, 200, 150)
            bbox2 = (300, 200, 180, 120)
            draw_bounding_box(image, bbox1, (0, 255, 0))
            draw_bounding_box(image, bbox2, (0,
255, 0)) distance =
            calculate_distance(bbox1, bbox2)
            if distance < 200:
                color = (0, 255, 0)
                print("Social Distancing")
            else:
                color = (0, 0, 255)
                print("Not maintaining Social Distancing")
            bbox_combined = (min(bbox1[0], bbox2[0]), min(bbox1[1], bbox2[1]),
                max(bbox1[0] + bbox1[2], bbox2[0] + bbox2[2]) - min(bbox1[0],
                bbox2[0]),
                max(bbox1[1] + bbox1[3], bbox2[1] + bbox2[3]) - min(bbox1[1],
                bbox2[1])) draw_bounding_box(image, bbox_combined, color)
            cv2.putText(image, f'Distance: {distance:.2f} pixels', (50, 30), cv2.FONT_HERSHEY_SIMPLEX, 1, (255, 255,
255), 2) cv2_imshow(image)

```

```

!pip install "openvino>=2024.0.0" "ultralytics==8.2.18" "torch>=2.1" "ipywidgets==7.7.1" from
pathlib import Path

from ultralytics import YOLO

models_dir = Path("./models")

models_dir.mkdir(exist_ok=True)

DET_MODEL_NAME =

"yolov8n"

det_model = YOLO(models_dir / f"{DET_MODEL_NAME}.pt")

label_map = det_model.model.names

res = det_model()

det_model_path=models_dir/f"{DET_MODEL_NAME}_openvino_model/{DET_MODEL_NAME}.x
ml" if not det_model_path.exists():

    det_model.export(format="openvino", dynamic=True, half=True)

from ultralytics import YOLO, solutions

import cv2 import

time import

collections

import numpy as np

from IPython import display

import torch

import openvino as ov

import ipywidgets as widgets

def run_inference(source, device):

    core = ov.Core()

    det_ov_model = core.read_model(det_model_path)

    ov_config = {}

    if "GPU" in device.value or ("AUTO" in device.value and "GPU" in core.available_devices):

        ov_config = {"GPU_DISABLE_WINOGRAD_CONVOLUTION": "YES"}

    compiled_model = core.compile_model(det_ov_model, device.value, ov_config)

    def infer(*args):

        result = compiled_model(args)

        return torch.from_numpy(result[0])

    det_model.predictor.inference = infer

    det_model.predictor.model.pt = False

    try:

        cap = cv2.VideoCapture(source)

        assert cap.isOpened(), "Error reading video file"

        line_points = [(0, 300), (1080, 300)]

        classes_to_count = [0]

        counter = solutions.ObjectCounter(view_img=False,reg_pts=line_points,

                                         classes_names=det_model.names,draw_tracks=True,

                                         line_thickness=2,view_in_counts=False,view_out_counts=False)

```

```

processing_times = collections.deque(maxlen=200)

while cap.isOpened():
    success, frame =
    cap.read() if not success:
        print("Video frame is empty or video processing has been successfully completed.")
        break
    start_time = time.time()
    tracks = det_model.track(frame, persist=True, show=False,
                             classes=classes_to_count, verbose=False)
    frame = counter.start_counting(frame, tracks)
    stop_time = time.time()
    processing_times.append(stop_time - start_time)
    _, f_width = frame.shape[:2]
    processing_time = np.mean(processing_times) *
    1000 fps = 1000 / processing_time
    cv2.putText(img=frame, text=f"Inference time: {processing_time:.1f}ms ({fps:.1f} FPS)",
               org=(20, 40), fontFace=cv2.FONT_HERSHEY_COMPLEX, fontScale=f_width / 1000, color=(0, 0, 255),
               thickness=2, lineType=cv2.LINE_AA)
    counts = counter.out_counts
    text = f"Count: {counts}"
    fontFace = cv2.FONT_HERSHEY_COMPLEX
    fontScale = 0.75
    thickness = 2
    (text_width, text_height), _ = cv2.getTextSize(text, fontFace, fontScale, thickness)
    top_right_corner = (frame.shape[1] - text_width - 20, 40)
    cv2.putText(img=frame, text=text, org=(top_right_corner[0],
    top_right_corner[1]), fontFace=fontFace, fontScale=fontScale, color=(0, 0, 255), thickness=thickness, lineType=cv2.LINE_AA)
    _, encoded_img = cv2.imencode(ext=".jpg", img=frame, params=[cv2.IMWRITE_JPEG_QUALITY,
    100]) i = display.Image(data=encoded_img)

```

```

import cv2
import numpy as np
from time import sleep
largura_min=80 #Largura minima do retangulo
altura_min=80 #Altura minima do retangulo
offset=6 #Erro permitido entre pixel
pos_linha=550 #Posição da linha de contagem
delay= 60 #FPS do vídeo
detec = []
carros= 0
def pega_centro(x, y, w, h):
    x1 = int(w / 2)
    y1 = int(h / 2)
    cx = x + x1
    cy = y + y1
    return cx,cy
cap = cv2.VideoCapture('video.mp4')
subtracao = cv2.bgsegm.createBackgroundSubtractorMOG()
while True:
    ret , frame1 = cap.read()
    tempo = float(1/delay)
    sleep(tempo)
    grey = cv2.cvtColor(frame1,cv2.COLOR_BGR2GRAY)
    blur = cv2.GaussianBlur(grey,(3,3),5)
    img_sub = subtracao.apply(blur)
    dilat = cv2.dilate(img_sub,np.ones((5,5)))
    kernel = cv2.getStructuringElement(cv2.MORPH_ELLIPSE, (5, 5))
    dilatada = cv2.morphologyEx (dilat, cv2. MORPH_CLOSE , kernel)
    dilatada = cv2.morphologyEx (dilatada, cv2. MORPH_CLOSE , kernel)
    contorno,h=cv2.findContours(dilatada,cv2.RETR_TREE,cv2.CHAIN_APPROX_SIMPLE)
    cv2.line(frame1, (25, pos_linha), (1200, pos_linha), (255,127,0), 3)
    for(i,c) in enumerate(contorno):
        (x,y,w,h) = cv2.boundingRect(c)
        validar_contorno = (w >= largura_min) and (h >= altura_min)
        continue
        cv2.rectangle(frame1,(x,y),(x+w,y+h),(0,255,0),2)
        centro = pega_centro(x, y, w, h)
        detec.append(centro)
        cv2.circle(frame1, centro, 4, (0, 0,255), -1)
    for (x,y) in detec:
        if y<(pos_linha+offset) and y>(pos_linha-offset):
            carros+=1
            cv2.line(frame1, (25, pos_linha), (1200, pos_linha), (0,127,255), 3)
            detec.remove((x,y))
            print("car is detected : "+str(carros))
            cv2.putText(frame1, "VEHICLE COUNT : "+str(carros), (450, 70),cv2.FONT_HERSHEY_SIMPLEX, 2, (0, 0, 255),5)
    cv2.imshow("Video Original" , frame1)
    cv2.imshow("Detector",dilatada)
    if cv2.waitKey(1) == 27:
        break
    cv2.destroyAllWindows()
    cap.release()
    cap.release()

```



```
display.clear_output(wait=True) display.display(i)
except KeyboardInterrupt:
    print("Interrupted") cap.release()
cv2.destroyAllWindows()
VIDEO_SOURCE = "https://github.com/intel-iot-devkit/sample-videos/raw/master/people-detection.mp4" import
ipywidgets as widgets
import opencv as cv core =
cv.Core()
device = widgets.Dropdown(options=core.available_devices + ["AUTO"], value="AUTO", description="Device:",
disabled=False)
device
run_inference(source=VIDEO_SOURCE, device = device)
```